

ARCHITECTURE · THE OUTCOME

The rigor of rules. The reach of AI.

A deterministic fraud engine an auditor can trust, wrapped by AI that reads the mess, writes the paperwork, and helps the rulebook grow — always with a human in the loop.

Deterministic core


AI at the edges

Human-approved rules

Runs offline

THE BIG IDEA


Rules alone are rigid. AI alone is a black box.



The mixture is the shot: a rulebook a regulator trusts, made **smarter and far less manual** by AI — without ever handing the AI the verdict.

SYSTEM AT A GLANCE

One claim in → a scored, explained, routed decision out



Purple = AI on the **edges** · green = the deterministic core the AI never decides. The Rule Lab loop lets the engine **grow**, human-gated.

THE “PROVABLY RIGHT” HALF

Every score is an itemised receipt

 Claim C-101 · auto_theft why the score is exactly what it is		
R1 · just-in-time cover		+15
R12 · shared phone + repairer		+15
R8 · threshold gaming		+12
R13 · flagged repairer		+10
R7 · above 90th-pct amount		+10
R3 · address mismatch		+8
R11 · no police report		+7
R6 · weak-witness timing		+5
TOTAL	82	HIGH → SIU

 **REPRODUCIBLE**
same input → same score

Why it matters: no probability to argue with — every point cites a named rule, and they simply **add up**. That's what makes a decline or an SIU referral **defensible**.

THE "ADAPTIVE" HALF

Three jobs for AI — none of them the verdict

 **READS**

 **messy text / PDF**

✨ structures

typed fields the rules read

 **WRITES**

 **scored claim / ring**

✨ writes

 **why-flagged note** ·  **ring summary** ·  **SIU referral**

 **EVOLVES**

 **your claim book**

✨ + 👤

 **new rule proposed** → human-adopted into core

In one line: AI turns **messy** → **structured**, **structured** → **readable**, and **data** → **better rules** — the number in the middle stays pure maths.

WHY IT COMPOUNDS

The rulebook gets smarter — **safely**



The moat: most AI fraud tools drift and can't explain themselves. Here the engine **improves with your book**, every change is a human-approved, auditable rule.

WHAT SINGLE-CLAIM REVIEW CAN'T SEE

Organised fraud = connected claims



In plain terms: three different people, but the same phone, shop and lawyer. That shared box tying strangers together is the ring — and it's the highest-value fraud to catch.

WHY IT PAYS

What it puts on the balance sheet

 \$0 AI running cost

Local **Ollama** on your own hardware — **no per-token API bills**, no cloud, and claim data **never leaves your network**. Add analysts & claims for free



↑ Detection

Catches organised rings & threshold-gaming a single-claim model structurally can't see.



↓ Leakage

Fraud runs **≈ 10% of claim payouts¹** — every miss is pure, un-recoverable loss.



Hours back

Notes, ring summaries & SIU referrals are **auto-drafted**, not hand-written.



Focus

A risk-ranked queue means adjusters work the **worst first** and skip the noise.



Defensible

Explainable + reproducible → fewer **fines** and overturned declines at audit.



Compounds

Rulebook grows with your book — **no ML retraining, no drift**, every change auditable.

¹ industry benchmark (Coalition Against Insurance Fraud). Figures illustrate the **value levers** this architecture unlocks — directional, not measured on the demo book.

THE ONE-LINER · THANK YOU

Team **Overwatch**

See the ring — not just the claim.

↑ catch

rings & hidden fraud

🕒 hours

paperwork auto-drafted

🛡️ audit

every decision defensible

📈 grows

with your data, safely

dashboard → `http://localhost:8000` · pitch → `/deck` · **103/103** tests green